


DevSecOps Platform - System Architecture

1. High-Level Architecture





2. Component Architecture

2.1 Core Services

A. Pipeline Orchestration Service

Purpose: Manage CI/CD pipelines, coordinate builds and deployments

Technologies: Go, Kubernetes API **Responsibilities:**

- Pipeline definition and execution
- Build orchestration
- Deployment management
- Rollback handling
- Environment configuration

APIs:

```
POST /api/v1/pipelines      # Create pipeline
GET /api/v1/pipelines/{id}   # Get pipeline status
POST /api/v1/pipelines/{id}/execute  # Trigger pipeline
POST /api/v1/pipelines/{id}/rollback # Rollback deployment
GET /api/v1/pipelines/history    # Pipeline history
```

B. Security Scanning Service

Purpose: Aggregate and manage security scan results

Technologies: Python (FastAPI), PostgreSQL **Responsibilities:**

- Coordinate SAST, SCA, DAST scans
- Vulnerability deduplication
- Risk scoring and prioritization
- Remediation tracking
- Compliance reporting

Integration Points:

- SonarQube API
- Semgrep CLI
- Gitleaks CLI
- Trivy API
- OWASP Dependency Check
- Checkov CLI
- OWASP ZAP API

APIs:

```

POST /api/v1/scans/sast      # Trigger SAST scan
POST /api/v1/scans/sca      # Trigger SCA scan
POST /api/v1/scans/dast      # Trigger DAST scan
GET /api/v1/vulnerabilities  # List vulnerabilities
GET /api/v1/vulnerabilities/{id} # Get vulnerability details
PUT /api/v1/vulnerabilities/{id} # Update vulnerability status
GET /api/v1/reports/security   # Generate security report

```

C. Testing Service

Purpose: Manage automated testing lifecycle

Technologies: Node.js (Express), MongoDB **Responsibilities:**

- Test execution orchestration
- Test result aggregation
- Performance testing
- Test reporting
- Flaky test detection

Integration Points:

- Selenium Grid
- Appium Server
- JMeter Cluster
- BrowserStack/Sauce Labs

APIs:

```
POST /api/v1/tests/functional      # Execute functional tests  
POST /api/v1/tests/performance    # Execute performance tests  
GET  /api/v1/tests/results        # Get test results  
GET  /api/v1/tests/coverage       # Get code coverage  
POST /api/v1/tests/suites         # Create test suite
```

D. IaC Security Service

Purpose: Validate infrastructure code security

Technologies: Python, PostgreSQL **Responsibilities:**

- Terraform plan analysis
- Kubernetes manifest validation
- Policy enforcement
- Compliance checking
- Remediation suggestions

APIs:

```
POST /api/v1/iac/scan            # Scan IaC files  
GET  /api/v1/iac/policies        # List policies  
POST /api/v1/iac/policies        # Create policy  
GET  /api/v1/iac/compliance     # Compliance status
```

E. HR Analytics Service

Purpose: Collect and analyze developer productivity metrics

Technologies: Python (FastAPI), TimescaleDB, PostgreSQL **Responsibilities:**

- Metrics collection from multiple sources
- Productivity calculation
- Team performance analysis
- Report generation
- Trend analysis

Data Sources:

- GitHub (commits, PRs, reviews)
- Jenkins (builds, deployments)
- SonarQube (code quality)
- Jira/Linear (tickets, velocity)

APIs:

```

GET /api/v1/analytics/individual/{userId}    # Individual metrics
GET /api/v1/analytics/team/{teamId}          # Team metrics
GET /api/v1/analytics/productivity           # Productivity trends
GET /api/v1/analytics/quality                # Quality metrics
POST /api/v1/analytics/reports               # Generate report

```

F. Workflow Automation Service (n8n Integration)

Purpose: Custom workflow automation and integration

Technologies: Node.js, n8n **Responsibilities:**

- Workflow execution
- Custom integration logic
- Notification handling
- Data transformation
- Scheduled tasks

G. Monitoring & Observability Service

Purpose: Platform health and performance monitoring

Technologies: Go, Prometheus, Grafana **Responsibilities:**

- Metrics collection
- Log aggregation
- Alert management
- Dashboard generation
- Performance tracking

2.2 External System Integrations

GitHub Integration

```
python

# GitHub Webhook Handler
- Events: push, pull_request, release, workflow_run
- Actions:
  - Trigger CI/CD on push
  - Scan code on PR
  - Update status checks
  - Comment on PRs with scan results
```

Jenkins Integration

```
python

# Jenkins Pipeline Triggers
- Job creation and configuration
- Build triggers
- Status monitoring
- Artifact management
- Build log collection
```

SonarQube Integration

```
python

# Code Quality Analysis
- Project creation
- Quality gate configuration
- Scan execution
- Issue retrieval
- Metrics collection
```

3. Data Models

3.1 Pipeline Schema (PostgreSQL)

sql

```
CREATE TABLE pipelines (
    id UUID PRIMARY KEY,
    name VARCHAR(255) NOT NULL,
    repository_url VARCHAR(500),
    branch VARCHAR(100),
    status VARCHAR(50),
    created_at TIMESTAMP,
    updated_at TIMESTAMP,
    created_by UUID,
    config JSONB
);
```

```
CREATE TABLE pipeline_executions (
    id UUID PRIMARY KEY,
    pipeline_id UUID REFERENCES pipelines(id),
    status VARCHAR(50),
    started_at TIMESTAMP,
    completed_at TIMESTAMP,
    duration INTEGER,
    triggered_by UUID,
    commit_sha VARCHAR(40),
    logs TEXT
);
```

```
CREATE TABLE pipeline_stages (
    id UUID PRIMARY KEY,
    execution_id UUID REFERENCES pipeline_executions(id),
    stage_name VARCHAR(100),
    status VARCHAR(50),
    started_at TIMESTAMP,
    completed_at TIMESTAMP,
    duration INTEGER
);
```

3.2 Security Schema (PostgreSQL)

sql

```
CREATE TABLE vulnerabilities (
    id UUID PRIMARY KEY,
    repository VARCHAR(255),
    branch VARCHAR(100),
    commit_sha VARCHAR(40),
    vulnerability_type VARCHAR(50), -- SAST, SCA, DAST, IAC
    severity VARCHAR(20), -- CRITICAL, HIGH, MEDIUM, LOW
    title VARCHAR(500),
    description TEXT,
    file_path VARCHAR(1000),
    line_number INTEGER,
    cve_id VARCHAR(50),
    cvss_score DECIMAL(3,1),
    status VARCHAR(50), -- OPEN, IN_PROGRESS, RESOLVED, FALSE_POSITIVE
    first_detected TIMESTAMP,
    last_detected TIMESTAMP,
    resolved_at TIMESTAMP,
    resolved_by UUID,
    metadata JSONB
);
```

```
CREATE TABLE scan_results (
    id UUID PRIMARY KEY,
    scan_type VARCHAR(50),
    repository VARCHAR(255),
    branch VARCHAR(100),
    commit_sha VARCHAR(40),
    started_at TIMESTAMP,
    completed_at TIMESTAMP,
    status VARCHAR(50),
    total_findings INTEGER,
    critical_count INTEGER,
    high_count INTEGER,
    medium_count INTEGER,
    low_count INTEGER,
    tool_name VARCHAR(100),
    tool_version VARCHAR(50)
);
```

3.3 Testing Schema (MongoDB)

javascript

```
{  
  "_id": ObjectId,  
  "test_suite_id": String,  
  "execution_id": String,  
  "test_type": String, //functional, performance, security  
  "status": String,  
  "started_at": ISODate,  
  "completed_at": ISODate,  
  "duration_ms": Number,  
  "results": {  
    "total": Number,  
    "passed": Number,  
    "failed": Number,  
    "skipped": Number  
  },  
  "test_cases": [{  
    "name": String,  
    "status": String,  
    "duration_ms": Number,  
    "error_message": String,  
    "stack_trace": String,  
    "screenshot_url": String,  
    "video_url": String  
  }],  
  "coverage": {  
    "line": Number,  
    "branch": Number,  
    "function": Number,  
    "statement": Number  
  },  
  "performance_metrics": {  
    "response_time_p50": Number,  
    "response_time_p95": Number,  
    "response_time_p99": Number,  
    "throughput": Number,  
    "error_rate": Number  
  }  
}
```

3.4 HR Analytics Schema (TimescaleDB)

sql

```
CREATE TABLE developer_metrics (
    timestamp TIMESTAMPTZ NOT NULL,
    user_id UUID NOT NULL,
    team_id UUID,
    metric_type VARCHAR(50), -- commits, prs, reviews, builds
    metric_value NUMERIC,
    metadata JSONB
);
```

```
SELECT create_hypertable('developer_metrics', 'timestamp');
```

```
CREATE TABLE productivity_snapshots (
    id UUID PRIMARY KEY,
    user_id UUID NOT NULL,
    snapshot_date DATE,
    commits_count INTEGER,
    lines_added INTEGER,
    lines_deleted INTEGER,
    prs_created INTEGER,
    prs_reviewed INTEGER,
    prs_merged INTEGER,
    code_review_time_avg INTEGER, -- minutes
    build_success_rate DECIMAL(5,2),
    bugs_introduced INTEGER,
    bugs_fixed INTEGER,
    code_quality_score DECIMAL(5,2),
    test_coverage DECIMAL(5,2)
);
```

4. Event Streaming Architecture (Kafka)

4.1 Kafka Topics

```
# Pipeline Events
dcc.pipeline.started
dcc.pipeline.stage.completed
dcc.pipeline.completed
dcc.pipeline.failed

# Security Events
dcc.security.scan.started
dcc.security.scan.completed
dcc.security.vulnerability.detected
dcc.security.vulnerability.resolved

# Testing Events
dcc.test.started
dcc.test.completed
dcc.test.failed

# HR Metrics Events
dcc.metrics.commit
dcc.metrics.pr
dcc.metrics.review
dcc.metrics.build

# Notification Events
dcc.notification.alert
dcc.notification.report
```

4.2 Event Schema Example

```
json
```

```
{
  "event_id": "uuid",
  "event_type": "dec.security.vulnerability.detected",
  "timestamp": "2025-10-15T10:30:00Z",
  "source": "security-scanning-service",
  "version": "1.0",
  "data": {
    "vulnerability_id": "uuid",
    "repository": "org/repo",
    "severity": "HIGH",
    "type": "SQL_INJECTION",
    "file": "src/api/user.py",
    "line": 42
  },
  "metadata": {
    "correlation_id": "uuid",
    "user_id": "uuid",
    "environment": "production"
  }
}
```

5. API Gateway Configuration

5.1 Kong Configuration

```
yaml
```

```
services:  
  - name: pipeline-service  
    url: http://pipeline-service:8080  
  routes:  
    - name: pipeline-routes  
      paths:  
        - /api/v1/pipelines  
      methods:  
        - GET  
        - POST  
        - PUT  
        - DELETE  
  plugins:  
    - name: jwt  
    - name: rate-limiting  
      config:  
        minute: 100  
    - name: correlation-id  
    - name: request-transformer  
  
  - name: security-service  
    url: http://security-service:8080  
  routes:  
    - name: security-routes  
      paths:  
        - /api/v1/scans  
        - /api/v1/vulnerabilities  
  plugins:  
    - name: jwt  
    - name: rate-limiting
```

6. Security Architecture

6.1 Authentication Flow

User → Admin Portal → API Gateway → Auth Service → JWT Token



Verify Token



Extract Claims



Check Permissions



Route to Service

6.2 RBAC Model

```
yaml
```

Roles:

```
- admin:
```

permissions:

```
  - "*"
```

```
- developer:
```

permissions:

- read:pipelines
- create:pipelines
- read:security
- update:vulnerabilities
- read:tests

```
- security_engineer:
```

permissions:

- "*:security"
- "*:vulnerabilities"
- read:pipelines

```
- manager:
```

permissions:

- read:*
- read:analytics
- read:reports

```
- hr_manager:
```

permissions:

- read:analytics
- read:hr_metrics
- create:reports

7. Deployment Architecture

7.1 Kubernetes Cluster Layout

Namespace: dcc-platform

- └─ api-gateway (Kong/Traefik)
- └─ auth-service
- └─ pipeline-service
- └─ security-service
- └─ testing-service
- └─ iac-service
- └─ hr-analytics-service
- └─ workflow-service
- └─ admin-portal-frontend
- └─ monitoring-stack

Namespace: dcc-integrations

- └─ sonarqube
- └─ jenkins
- └─ n8n
- └─ selenium-grid
- └─ kafka-cluster
- └─ redis


Namespace: dcc-data

- └─ postgresql
- └─ mongodb
- └─ timescaledb
- └─ elasticsearch

Namespace: dcc-monitoring

- └─ prometheus
- └─ grafana
- └─ logstash
- └─ kibana
- └─ jaeger

7.2 High Availability Setup



8. Monitoring & Observability

8.1 Metrics Collection

Application Metrics (Prometheus):

- Request rate, latency, errors (RED method)
- Resource utilization (CPU, Memory)
- Custom business metrics

Infrastructure Metrics:

- Node health
- Pod status
- Network throughput

Database Metrics:

- Query performance
- Connection pool
- Replication lag

8.2 Logging Strategy

ELK Stack Pipeline:

Application Logs → Filebeat → Logstash → Elasticsearch → Kibana

Log Levels:

- ERROR: Critical issues requiring immediate attention
- WARN: Potential issues
- INFO: Key business events
- DEBUG: Detailed debugging information

Structured Logging Format (JSON):

```
{  
  "timestamp": "2025-10-15T10:30:00Z",  
  "level": "INFO",  
  "service": "pipeline-service",  
  "message": "Pipeline execution started",  
  "correlation_id": "uuid",  
  "user_id": "uuid",  
  "metadata": {}  
}
```

8.3 Distributed Tracing

Jaeger Integration:

- Track request flow across microservices
- Identify performance bottlenecks
- Debug complex interactions
- Visualize service dependencies

9. Disaster Recovery

9.1 Backup Strategy

Databases:

- PostgreSQL: Daily full backup, hourly incremental
- MongoDB: Daily backup with point-in-time recovery
- TimescaleDB: Continuous archiving

Configuration:

- GitOps approach: All configs in Git
- Secrets: Encrypted backups to secure storage

Recovery Time Objective (RTO): 4 hours

Recovery Point Objective (RPO): 1 hour

9.2 Failover Procedures

1. Automatic Pod Restart (K8s Health Checks)
2. Node Failure → Pod Migration
3. Database Failover → Promote Replica
4. Region Failure → DNS Failover to Secondary Region

10. Performance Optimization

10.1 Caching Strategy

Redis Layers:

- L1: API Response Cache (TTL: 5 minutes)
- L2: Database Query Cache (TTL: 15 minutes)
- L3: Session Cache

Cache Invalidation:

- Event-driven invalidation via Kafka
- TTL-based expiration
- Manual purge via API

10.2 Database Optimization

PostgreSQL:

- Connection pooling (PgBouncer)
- Partitioning for large tables
- Indexing strategy
- Query optimization

MongoDB:

- Sharding for horizontal scaling
- Compound indexes
- Read replicas for analytics

TimescaleDB:

- Hypertables for time-series data
- Compression policies
- Retention policies

11. Scalability Considerations

11.1 Horizontal Scaling

Services:

- Stateless design
- Horizontal Pod Autoscaler (HPA)
- Load balancing
- Session affinity when needed

Databases:

- Read replicas for read-heavy workloads
- Sharding for write-heavy workloads
- Connection pooling

Message Queue:

- Kafka partitions for parallel processing
- Consumer groups for load distribution

11.2 Capacity Planning

Baseline Metrics:

- 1000 concurrent users
- 100 builds per hour
- 10,000 scan results per day
- 1 million HR metrics per day

Scaling Triggers:

- CPU > 70% for 5 minutes → Scale up
- Request latency p95 > 500ms → Scale up
- Queue depth > 1000 → Scale up

This architecture provides a robust, scalable, and secure foundation for the DevSecOps platform.