



# Enterprise API Management Platform - Architecture Diagrams & Component Interactions

---

## 1. REQUEST FLOW DIAGRAM (End-to-End)





## UPSTREAM SERVICE (Backend API)

- Process request
- Return response (200, 400, 500, etc.)

4. Response from Upstream

## KONG DATA PLANE (Response Processing)

### Response Processing Pipeline

#### 1. Status Code Validation

- └ Log response status

#### 2. Response Transformation

- └ Add security headers (HSTS, CSP, etc.)
- └ Remove sensitive headers
- └ Transform body if needed
- └ Set cache-control headers

#### 3. PII Masking

- └ Detect & redact sensitive data
- └ Log masked version to audit trail

#### 4. Compression


- └ gzip or brotli compress
- └ Add Content-Encoding header
- └ Reduce response size

#### 5. Metrics Recording

- └ Stop request timer
- └ Calculate latency percentiles
- └ Record status code distribution
- └ Update error counters if needed
- └ Export to Prometheus /metrics endpoint

#### 6. Audit Logging (Async)

- └ Prepare structured log entry
- └ Include: timestamp, tenant, user, path, method, status, latency
- └ Send to Kafka topic (async, buffered)



#### BACKGROUND ASYNC OPERATIONS (Non-blocking):

- Rate limit counter sync to distributed store (every 10-30 sec)
- Workflow trigger event processing (NATS → Temporal)
- Audit log batch write to Kafka (every 5-10 sec)
- Metrics aggregation and export (every 10 sec)
- Config cache refresh from control plane (every 30 sec)

## 2. ARCHITECTURE COMPONENT INTERACTIONS

## KONG DATA PLANE (Stateless, Horizontal Scaling)

REQUEST → Authentication → Authorization → Rate Limiting →  
Transformation → Logging → Upstream → Response

↓ Dependencies:

- └ Control Plane (Pull config every 30 sec)
- └ Vault (Fetch API keys/credentials, cache 1-5 min)
- └ Redis (Distributed rate limit state, TTL 1 min)
- └ Kafka (Send audit logs async)
- └ NATS (Receive config invalidation, send workflow triggers)
- └ Upstream Services (Route requests)
  - └ Circuit breaker: Fail fast if service down
- └ (State: Last failure timestamp, trip status)

## KONG CONTROL PLANE

Admin API:

- └ POST /services (Create API service)
- └ POST /routes (Create API endpoint)
- └ POST /plugins (Configure plugins)
- └ GET /config (Return current state)
- └ DELETE / PATCH (Modify entities)

↓ Storage:

- └ PostgreSQL HA (Primary + 2-3 replicas)
  - └ Synchronous replication (zero RPO)
  - └ Automatic failover via Patroni

↓ Integration:

- └ Data Plane: Pull changes (pull-based, no push load)
- └ Vault: Store credentials, fetch as needed
- └ Kafka: Publish config change events
- └ NATS: Broadcast invalidation signals

## TEMPORAL WORKFLOW ENGINE

Workflow: Long-running business process

Activity: Individual step (can fail & retry)

Examples:

Workflow: Process payment (validate → charge → notify)

Activity: Call payment API (can timeout/fail)

Automatic Retry: Exponential backoff up to max attempts

↓ Components:

Frontend: gRPC endpoint for workflow clients

History Service: Event store in PostgreSQL

Matching Service: Task queue management

Worker: Execute activities (custom code)

Visibility Store: Query workflow state

↓ Integration:


Kong: Receive trigger events (API call → workflow)

NATS: Listen for signals (pause/resume/cancel)

PostgreSQL: Persist workflow state

Kafka: Log all workflow events (audit)

Jaeger: Send traces for distributed tracing



## NATS JETSTREAM (Local Per-Region, Low-Latency)

### Subjects (Stream Subjects):

└ workflow.triggers.{tenant\_id}

  └ Producer: Kong when API matches trigger rule

  └ Consumer: Temporal (subscribe to trigger events)

  └ Latency: Sub-millisecond

└ workflow.signals.{tenant\_id}

  └ Producer: External systems, APIs

  └ Consumer: Temporal (send signal to workflow)

  └ Use case: Approve/reject workflow manually

└ config.invalidations

  └ Producer: Kong control plane (config changed)

  └ Consumer: All Kong data planes (broadcast)

  └ Action: Purge local cache, fetch fresh config

  └ Delivery: Exactly-once guarantees

└ health.heartbeats

  └ Producer: All services (every 30 sec)

  └ Consumer: Monitoring/alerting system

  └ Action: Detect service failures in <10 sec

### ↓ Multi-Tenancy:

  └ Subjects scoped per tenant → No cross-tenant leakage

  └ Subjects: workflow.triggers.tenant-123

  └ Authorization: JWT-based per-tenant

  └ Isolation: Complete (no way to access other tenants)

## VAULT (Secrets Management)

### Secret Types:

#### API Keys (for 3rd-party services)

Path: secret/data/3rdparty/{service}/{api\_key}

Rotation: Every 30-90 days

Accessed by: Kong plugins (via HTTP API)

Cache: 5 minutes in-memory

#### OAuth Secrets

Path: secret/data/oauth/{provider}/{client\_secret}

Accessed by: OAuth2 plugin

#### mTLS Certificates

Path: secret/data/certs/{service}/{cert\_pem}

Rotation: Automatic via cert-manager (1 week before)

Cached: In Kong worker memory

#### Database Credentials

Path: secret/data/db/{environment}/{role}

Rotation: Every 30 days

### ↓ High Availability:

#### 5-node HA cluster (Raft consensus)

#### Automatic failover

#### Encrypted storage (KMS + encryption at rest)

#### Backup: Daily snapshots to S3 (encrypted)

### ↓ Audit:

#### All secret access logged with: who, what, when, result

## POSTGRESQL (Operational State Store)

### Tables:

#### workspaces (Tenant definitions)

id, name, description, created\_at

Used for multi-tenancy isolation

#### rbac\_roles (Permission definitions)

admin, operator, developer, viewer, custom

Assigned to users per workspace

#### services (API backend definitions)

name, url, protocol, health\_check\_url

Partitioned by workspace\_id

Referenced by routes

#### routes (API endpoints exposed to clients)

name, service\_id, path, methods, hosts

upstream\_url, transformations

Indexed by: workspace\_id, path, method

Returns by GET /api/routes?workspace=X

#### plugins (Functionality extensions)

name (auth, ratelimit, transform, etc.)

service\_id or route\_id (applies to which entity)

config (JSON) - plugin-specific settings

enabled (boolean)

#### rate\_limiting\_policies (Traffic control)

workspace\_id, route\_id, user\_id (optional)

limit (e.g., 1000 requests)

window (e.g., 1 hour)

algorithm (sliding window, token bucket)

Cached in Redis for performance

#### audit\_logs (Immutable change log)

timestamp, user\_id, action, entity\_type, entity\_id

before, after (JSON snapshots)

Partitioned by (workspace\_id, date)

Immutable: INSERT only, no DELETE

TTL-based purging (default 2 years)

```
api_calls (Request audit trail)
  timestamp, tenant_id, user_id, method, path
  status_code, latency_ms, upstream_url
  error_message (if any), response_size
  Partitioned by (tenant_id, date)
  Sampled logging (100% errors, 1% success)
  TTL-based purging (default 90 days)

workflow_triggers (Event → Workflow mappings)
  workspace_id, route_id, event_type
  workflow_name, workflow_input (JSON template)
  enabled (boolean)
  Created via UI/API

↓ High Availability:
  Primary + 2-3 Replicas
  Synchronous replication (zero RPO)
  Automatic failover via Patroni
  WAL archiving to S3 (point-in-time recovery)
  Daily snapshots (backed up to S3)
```

### 3. MULTI-TENANT ISOLATION FLOW

REQUEST ARRIVES WITH TENANT CONTEXT

TENANT EXTRACTION (Priority Order)

1. X-Tenant-ID Header

└ Explicit tenant ID from client

2. JWT Claims

└ Extract tenant from JWT token (exp: tenant\_id)

3. Host Header (Subdomain)

└ {tenant-id}.api.example.com → tenant-id

4. API Key Lookup

└ Query DB: API key → workspace\_id

5. mTLS Certificate Subject

└ Extract tenant from cert CN or SAN

STORE TENANT IN REQUEST CONTEXT

context.tenant\_id = "tenant-123"

context.workspace\_id = 42

ALL DOWNSTREAM OPERATIONS USE THIS CONTEXT

DATA PLANE OPERATIONS (All Tenant-Scoped)

1. Authentication

└ Validate creds belong to tenant

2. Authorization

└ Query RBAC: user\_id + tenant\_id → permissions

```
3. Route Resolution
└ SELECT * FROM routes WHERE workspace_id = 42

4. Rate Limiting
└ KEY = "tenant-123:user-id:endpoint" in Redis

5. Logging
└ Write to Kafka: {tenant_id: "tenant-123", ...}
└ Partition by tenant_id (co-locate events)
```

#### RESPONSE PHASE (Tenant Context Preserved)

```
1. Add Response Headers
└ X-Tenant-ID: tenant-123 (for client audit)

2. Log Response
└ Include tenant_id in all logs

3. Send to Kafka
└ Partition: tenant-123 (keeps logs together)
└ Consumer sees only tenant's logs

4. Trigger Workflow (If Applicable)
└ NATS subject: workflow.triggers.tenant-123
  (tenant-scoped visibility)
```

#### KEY ISOLATION PRINCIPLES:

- └ Tenant ID in all queries (no global scope access)
- └ Kafka partitioned by tenant\_id (co-location)
- └ NATS subjects scoped per tenant (ACL-enforced)
- └ RBAC enforces cross-tenant access denial
- └ Row-level security in DB (tenant\_id filter)
- └ Audit logs capture all cross-boundary attempts

## 4. WORKFLOW INTEGRATION FLOW

API REQUEST ARRIVES AT KONG

WORKFLOW TRIGGER PLUGIN (In Kong Chain)

```
if route.id matches workflow_trigger rule:  
{  
  workflow_name: "process_payment"  
  tenant_id: extract_from_request()  
  input: {  
    user_id: request.user_id,  
    amount: request.body.amount,  
    currency: request.body.currency  
  }  
}
```

TRIGGER EVENT PUBLISHED TO NATS

(Async, non-blocking)

```
Subject: workflow.triggers.tenant-123  
Message: {  
  workflow_name: "process_payment",  
  workflow_id: "uuid-generated",  
  tenant_id: "tenant-123",  
  input: {...},  
  api_call_id: "request-id-from-header"  
}
```

ADD TO RESPONSE HEADERS

```
X-Workflow-Execution-ID: uuid-generated  
X-Workflow-Status-Endpoint: /workflow/status/{uuid}
```

Client can poll for workflow status

REQUEST CONTINUES TO UPSTREAM (Doesn't Wait)

Response is returned to client

Workflow executes in parallel

▼ (Async, parallel execution)

KONG/CLIENT

CONTINUES

Returns 200

Response

TEMPORAL WORKFLOW EXECUTION

(Separate thread/process)

Workflow: process\_payment

Steps:

1. Validate input (Activity)

2. Call payment gateway (Activity)

3. Log transaction (Activity)

4. Send confirmation (Activity)

Automatic Retries:

Transient failures: Exp backoff

Service unavailable: Retry up to 3x

Permanent failures: Alert operator

State Persistence:

Events logged to PostgreSQL

(can replay from last known state)

Audit Trail:

Each activity logged to Kafka  
(workflow.events topic)

↓

**WORKFLOW COMPLETION**

  |

  | Update state in Temporal

  | Log final result to DB

  | Send notification event

  | (via NATS or Kafka)

**CLIENT CAN QUERY WORKFLOW STATUS:**

  |— GET /api/workflows/{execution\_id}

  |— Returns: current\_step, status, progress, last\_update

  |— Webhook callback available (when workflow completes)

---

## 5. THIRD-PARTY API INTEGRATION FLOW

INTERNAL SERVICE NEEDS TO CALL 3RD-PARTY API



SERVICE CALLS KONG OUTBOUND PROXY

```
POST /outbound/3rdparty/stripe/charges
{
  "internal_api_key": "xxx",
  "request_body": {...},
  "tenant_id": "tenant-123"
}
```



KONG OUTBOUND PROXY PLUGIN

1. AUTHENTICATE INTERNAL CALLER

- └ Verify internal\_api\_key
  - (can be service-to-service mTLS)

2. FETCH 3RD-PARTY CREDENTIALS FROM VAULT

- └ Service account: Kong process
- └ Path: secret/data/3rdparty/stripe/api\_key
- └ Cache: 5 min in memory

3. PREPARE REQUEST TO 3RD-PARTY

- └ Add Stripe API key to header
- └ Add correlation ID for tracing
- └ Add X-Tenant-ID for logging
- └ Validate request against schema



CIRCUIT BREAKER LOGIC

Check State:  
└ CLOSED: Normal operation, forward request

- OPEN: Service down, return cached error
- HALF\_OPEN: Testing if recovered

Thresholds:

- Failure rate > 50% for 30 sec → OPEN
- Timeout > 5 sec (3 consecutive) → OPEN
- Recovery test every 60 sec → HALF\_OPEN

▼ (If not open)

### CALL 3RD-PARTY SERVICE

- HTTP POST <https://api.stripe.com/v1/charges>
- Headers:
  - Authorization: Bearer {api\_key\_from\_vault}
  - X-Correlation-ID: {from\_request}
  - X-Tenant-ID: tenant-123
  - X-Request-Timeout: 5s
- With Timeout:
  - 5 second timeout (fail fast)

### RESPONSE HANDLING

- 2xx (Success):
  - Update circuit breaker (success)
  - Cache successful response (if applicable)
  - Log to Kafka topic: third-party.interactions

- 4xx (Client Error):
  - Return error to caller (bad request)
  - Log for debugging (alert if suspicious)
  - Circuit breaker: Don't penalize (client issue)

- 5xx (Service Error):
  - Increment failure counter
  - Check circuit breaker state
  - If failures > threshold → OPEN circuit

```
  └─ Retry logic (with exponential backoff)
    └─ Retry 1: Wait 100ms
    └─ Retry 2: Wait 200ms
    └─ Retry 3: Wait 400ms
  └─ If all retries fail → Return error
  └─ Timeout:
    └─ Treat as service error
    └─ Increment failure counter
    └─ Retry with backoff
    └─ Circuit breaker may open if too many
```

```
LOG TO KAFKA (Async)
  └─ Topic: third-party.interactions
  └─ Partition Key: tenant_id + service_name
  └─ Message:
    {
      timestamp: ISO8601,
      tenant_id: "tenant-123",
      service: "stripe",
      method: "POST",
      endpoint: "/v1/charges",
      request_id: "req-xxx",
      status_code: 200,
      latency_ms: 234,
      request_size: 512,
      response_size: 2048,
      error_code: null,
      circuit_breaker_state: "CLOSED"
    }
  └─ Consumers:
    └─ Analytics: Aggregate metrics per service
    └─ Monitoring: Alert on high failure rates
    └─ Billing: Track API usage for chargeback
    └─ Audit: Immutable record for compliance
```

```
↓  
| RETURN RESPONSE TO CALLER  
|  
| HTTP 200 OK  
| {  
|   "status": "success",  
|   "data": {...}, (3rd-party response)  
|   "correlation_id": "req-xxx"  
| }  
|  
| Or on failure:  
| HTTP 502 Bad Gateway / 503 Service Unavailable  
| {  
|   "error": "3rd-party service unavailable",  
|   "correlation_id": "req-xxx",  
|   "retry_after": 60 (seconds)  
| }  
|
```

## 6. SCALABILITY FLOW (Auto-Scaling in Action)

#### TIME: 10:00 AM - BASELINE LOAD

- └ Kong Data Plane: 3 nodes (1 per AZ)
- └ CPU usage: ~40% per node
- └ Memory: ~200MB per node
- └ Status: Healthy

#### TIME: 10:15 AM - TRAFFIC SPIKE DETECTED

- └ Requests/sec jump from 5K to 12K
- └ CPU usage rises to 75% on all nodes
- └ Memory usage increases to ~400MB
- └ HPA (Horizontal Pod Autoscaler) triggered:
  - └ Threshold: CPU > 70% for 30 sec
  - └ Action: Scale from 3 → 6 replicas

#### TIME: 10:18 AM - KUBERNETES SCALING

- └ Cluster Autoscaler checks available capacity
- └ Scenario 1: Capacity available on existing nodes
  - └ Deploy pods to free slots (1-2 sec)
- └ Scenario 2: No available capacity
  - └ Trigger cluster autoscaler:
    - └ Request new worker node from cloud
    - └ AWS/GCP: Provision takes 1-2 minutes
    - └ KEDA: Scale based on Kafka lag (if applicable)

#### TIME: 10:20 AM - STEADY STATE AT HIGH LOAD

- └ Kong Data Plane: 6 nodes (+ extras may be provisioning)
- └ CPU usage: ~45% per node (balanced across 6)
- └ Memory: ~250MB per node
- └ Requests distributed evenly (load balancer)
- └ Status: Healthy, fully scaled

#### TIME: 10:35 AM - TRAFFIC SUBSIDES

- └ Requests/sec drops back to 6K
- └ CPU usage: ~35% per node
- └ HPA scale-down: Enabled after 5 min of low CPU

- └ Action: Scale from 6 → 4 replicas
  - └ Wait for graceful termination (30 sec drain)
  - └ Existing requests complete
  - └ Pods drain connections

TIME: 10:40 AM - BACK TO BASELINE

- └ Kong Data Plane: 4 nodes (consolidating)
- └ Cluster Autoscaler: May remove excess node (5 min delay)
- └ CPU usage: ~40% per node
- └ Status: Normal

SCALING METRICS:

- └ CPU Utilization: 40-80% target range
- └ Memory Usage: 200-500MB per node
- └ Kafka Consumer Lag: <10 sec (for event processing)
- └ Request Latency (p99): <500ms
- └ Error Rate: <0.1%

AUTOSCALING CONFIGURATION:

- └ HPA Min Replicas: 3 (HA requirement)
- └ HPA Max Replicas: 50 (cost control)
- └ CPU Threshold: 70%
- └ Memory Threshold: 80%
- └ Scale-up delay: 30 sec
- └ Scale-down delay: 5 min
- └ KEDA enabled: YES (for Kafka lag-based scaling)
- └ Graceful termination period: 30 sec

## 7. OBSERVABILITY INTEGRATION MAP

## KONG DATA PLANE

(Every request passes through here)

Prometheus | OpenTelemetry | Kafka | Jaeger |

Metrics | Tracing | Logs | Traces |

/metrics | /traces | Topics | Collector |

## Observability Stack (Central Region)

### Prometheus Server (HA)

— Scrapes metrics every 15 sec

— Retention: 15 days hot

— Remote storage: Thanos (long-term)

— Queries: Grafana dashboards

### Jaeger (Distributed Tracing)

— Collector: Receives spans

— Query Service: UI for traces

— Storage: Elasticsearch

— Sampling: 1-10% + 100% on errors


### ELK Stack (Log Aggregation)

— Elasticsearch: Indexed storage

— Logstash: Log parsing

— Kibana: UI + dashboards

— Retention: 90 days hot, 2y cold



#### METRIC EXAMPLES (Exported by Kong):

- └─ http\_requests\_total{method="POST", path="/api/users", status="200", tenant="tenant-123"}
- └─ http\_request\_duration\_seconds{quantile="0.99", path="/api/users"}
- └─ kong\_upstream\_target\_health{upstream="backend-service", target="10.0.1.5:8080", state="up"}
- └─ rate\_limit\_check\_duration\_seconds (custom)
- └─ auth\_token\_validation\_duration\_seconds (custom)

#### TRACE EXAMPLES (Sent to Jaeger):

- └─ Trace: 4a2c461b1d6a3b5e
  - └─ Span 1: Kong Request Processing (10ms)
    - └─ Event: TLS Handshake (2ms)
    - └─ Event: Auth Plugin (1ms)
    - └─ Event: Rate Limit Check (0.5ms)
    - └─ Event: Route Resolution (0.1ms)
  - └─ Span 2: Upstream Service Call (45ms)
    - └─ Event: DNS Lookup (1ms) [if needed]
    - └─ Event: Connect (2ms)
    - └─ Event: Request/Response (42ms)
  - └─ Span 3: Kong Response Processing (5ms)
    - └─ Event: Transformation (2ms)
    - └─ Event: Logging (3ms)
- └─ Total Trace: 60ms (visible in Jaeger UI)

#### LOG EXAMPLES (Sent to ELK):

```
└─ {timestamp: "2025-10-22T14:30:15Z", tenant: "tenant-123", user: "user-456", method: "POST", path: "/api/payments", status: 200, latency_ms: 234, request_id: "req-abc"}  
└─ {timestamp: "2025-10-22T14:30:20Z", tenant: "tenant-123", error: "Rate limit exceeded", limit: 1000, window: "1h"}  
└─ {timestamp: "2025-10-22T14:30:25Z", tenant: "tenant-123", auth_error: "Invalid API key"}
```

#### ALERT EXAMPLES (Triggered from Prometheus):

```
└─ CRITICAL: Kong Data Plane Down  
  └─ Condition: No nodes responding for 1+ min  
  
└─ WARNING: High Error Rate  
  └─ Condition: Error rate > 1% for 5 min  
  
└─ CRITICAL: P99 Latency High  
  └─ Condition: p99 latency > 500ms for 5 min  
  
└─ WARNING: Kafka Lag High  
  └─ Condition: Consumer lag > 10K messages  
  
└─ INFO: Scaling Event  
  └─ Triggered: 6 nodes (up from 3)
```

## 8. FAILOVER & DISASTER RECOVERY FLOW

## SCENARIO 1: SINGLE DATA PLANE NODE FAILURE

- |- 10:15 AM: Node crashes (OOM, kernel panic, etc.)
- |- Detection (K8s)
  - |- Liveness probe fails (3 consecutive failures)
    - |- Latency: ~30 seconds
- |- Action (K8s)
  - |- Mark pod as "Not Ready"
  - |- Load balancer removes from target group
  - |- Pod evicted
  - |- Replica set spawns new pod on different node
- |- Recovery
  - |- New pod starts (~10 sec)
  - |- Rejoins load balancer (~5 sec)
  - |- RTO: ~45 seconds total
  - |- Data loss: ZERO (stateless data plane)
- |- Impact: Minimal (traffic distributed to other 4 nodes)

## SCENARIO 2: ENTIRE KUBERNETES CLUSTER FAILURE (Regional)

- |- 10:15 AM: Cluster networking failure
  - |- etcd cluster becomes unavailable
- |- Detection
  - |- Multiple services report high latency
  - |- Prometheus scrapers timeout
  - |- Alerting: Critical alert fires
- |- Action (Manual or Automatic)
  - |- Automatic: DNS failover to secondary region
  - |- Manual: Update Route 53 to secondary region
  - |- Latency: 30-60 sec (TTL propagation)
- |- Secondary Region Activation
  - |- Kong data planes in secondary region already running
  - |- PostgreSQL: Promote secondary replica to primary
    - |- Current replication lag: <100ms
  - |- Data loss: Zero to minimal
  - |- Promotion time: ~30 sec

- └ Temporal: Workflow state in replicated DB
  - └ Workflows resume from last persisted state
- └ Kafka: Multi-region replication (RF=2)
  - └ Secondary cluster already has events
- └ Recovery Steps
  1. Detect primary region unavailable (2 min)
  2. DNS failover to secondary (1 min)
  3. Promote secondary DB (1 min)
  4. Validate secondary region health (1 min)
  - └ Total RTO: ~5 minutes
- └ Data Consistency
  - └ In-flight requests: May be lost (acceptable)
  - └ Persisted state: Recovered from replicas
  - └ Workflows: Replay from last committed state
  - └ RPO: <5 minutes (sync replication to secondary)
- └ Post-Recovery
  - └ Investigate root cause (human)
  - └ Restore primary region (if applicable)
  - └ Failback to primary (planned)

### SCENARIO 3: POSTGRESQL PRIMARY FAILURE

- └ 10:15 AM: Primary DB becomes unavailable
- └ Detection (Patroni)
  - └ Replica loses connection to primary
  - └ Quorum broken (can't reach primary)
  - └ Latency: ~10 seconds
- └ Automatic Failover (Patroni)
  - └ Quorum elects best replica as new primary
  - └ Promotion: Secondary → Primary (with standby)
  - └ TTL: <30 seconds
  - └ Write operations resume on new primary
- └ Impact
  - └ Read operations: Minimal impact (secondary already handling)
  - └ Write operations: Brief pause (<30 sec)

- | └ In-flight transactions: May roll back
- | └ Data loss: Zero (synchronous replication)
- |
- | └ Post-Failover
  - | └ Verify replication to remaining secondaries
  - | └ Spin up new standby (from backup or rebuild)
  - | └ Update Route 53 (already updated by Patroni)

#### SCENARIO 4: KAFKA BROKER FAILURE

- |
- | └ 10:15 AM: Kafka broker 2/3 fails
- |
- | └ Kafka Cluster State
  - | └ Cluster: 3 brokers (RF=3 by default)
  - | └ Leader election: New broker elected
  - | └ Replication: Continues on healthy brokers
  - | └ Latency: ~5 seconds (Zookeeper/KRaft election)
- |
- | └ Impact
  - | └ Audit logs: Continue writing to other brokers
  - | └ Event stream: No data loss (RF=3)
  - | └ Producers: May see increased latency
  - | └ Consumers: No impact
- |
- | └ Recovery
  - | └ Identify failed broker
  - | └ Replace/restart broker
  - | └ Wait for replication catch-up (5-30 min depending on backlog)
  - | └ RTO: <5 min to regain full capacity
- |
- | └ Data Protection
  - | └ Backups: Daily snapshots to S3 (cross-region)
  - | └ Can restore entire cluster if needed

#### DISASTER RECOVERY PROCEDURES:

- |
- | └ Backup Strategy
  - | └ PostgreSQL: Daily snapshot to S3 (cross-region)
  - | └ Kafka: Replication RF=2-3 (no explicit backups needed)
  - | └ Vault: Daily snapshots encrypted to S3
  - | └ Recovery: Can restore to any point in time

- └ RPO (Recovery Point Objective)
  - └ PostgreSQL: <5 min (sync replication)
  - └ Kafka: Real-time (3 replicas)
  - └ Configuration: <30 sec (config caching)
  - └ Audit logs: <1 sec (buffered writes)
- └ RTO (Recovery Time Objective)
  - └ Data plane failure: <1 min (auto-recovery)
  - └ Regional failure: <5 min (DNS + promotion)
  - └ Database failure: <1 min (Patroni promotion)
  - └ Full cluster rebuild: <30 min (from backups)
- └ Testing & Validation
  - └ Quarterly: Run full DR simulation
  - └ Monthly: Test individual component failure
  - └ After deploy: Chaos engineering tests
  - └ Monthly: Backup restoration test

## **End of Architecture Diagrams Document**

This document provides visual flows and detailed component interactions to complement the main architectural documentation.