

NexusTrade SuperApp - Deployment Guide

Version: 1.0.0

Last Updated: October 14, 2025

Table of Contents

1. [Pre-Deployment Checklist](#)
 2. [Infrastructure Setup](#)
 3. [Database Configuration](#)
 4. [Service Deployment](#)
 5. [Network Configuration](#)
 6. [Security Hardening](#)
 7. [Monitoring Setup](#)
 8. [Post-Deployment Verification](#)
 9. [Rollback Procedures](#)
 10. [Troubleshooting](#)
-

Pre-Deployment Checklist

Infrastructure Requirements

- Kubernetes cluster (1.29+) with 5 PoPs configured
- Load balancers configured for each PoP
- DNS configured with Geo-DNS routing
- SSL/TLS certificates issued
- Secrets management (Vault) configured
- Container registry accessible
- CI/CD pipeline configured

Access & Credentials

- AWS/Azure/GCP credentials configured
- Kubernetes cluster access verified
- Container registry credentials set up
- Database master passwords generated
- API keys for third-party services obtained
- Vault unsealed and accessible

Dependencies

- Docker images built and pushed
 - Helm charts validated
 - Database schemas prepared
 - Configuration files reviewed
 - Environment variables documented
-

Infrastructure Setup

1. Provision Kubernetes Clusters

Option A: Using Terraform (Recommended)

```
bash

cd infrastructure/terraform/aws

# Initialize Terraform
terraform init

# Review plan
terraform plan -var-file=production.tfvars

# Apply configuration
terraform apply -var-file=production.tfvars

# Repeat for all 5 PoPs
```

Option B: Manual Provisioning

Lagos (Primary PoP):

```
bash

# Create EKS cluster
eksctl create cluster \
--name nexustrade-lagos \
--version 1.29 \
--region af-south-1 \
--nodegroup-name standard-workers \
--node-type t3.xlarge \
--nodes 3 \
--nodes-min 3 \
--nodes-max 10 \
--managed
```

Repeat for:

- Johannesburg: `(af-south-1)`
- London: `(eu-west-2)`
- Ashburn: `(us-east-1)`
- Singapore: `(ap-southeast-1)`

2. Configure kubectl Contexts

```
bash

# Add contexts for all PoPs
kubectl config use-context nexustrade-lagos
kubectl config use-context nexustrade-joburg
kubectl config use-context nexustrade-london
kubectl config use-context nexustrade-ashburn
kubectl config use-context nexustrade-singapore

# Verify access
kubectl get nodes --all-namespaces
```

3. Install Core Components

```
bash

# Install Helm
curl https://raw.githubusercontent.com/helm/helm/main/scripts/get-helm-3 | bash

# Add Helm repositories
helm repo add prometheus-community https://prometheus-community.github.io/helm-charts
helm repo add grafana https://grafana.github.io/helm-charts
helm repo add bitnami https://charts.bitnami.com/bitnami
helm repo update

# Install cert-manager
kubectl apply -f https://github.com/cert-manager/cert-manager/releases/download/v1.13.0/cert-manager.yaml

# Install ingress-nginx
helm install ingress-nginx ingress-nginx/ingress-nginx \
--namespace ingress-nginx \
--create-namespace \
--set controller.service.type=LoadBalancer
```

Database Configuration

1. Deploy YugabyteDB Cluster

```
bash

# Create namespace
kubectl create namespace nexustrade

# Add YugabyteDB Helm repo
helm repo add YugabyteDB https://charts.yugabyte.com

# Install YugabyteDB
helm install YugabyteDB YugabyteDB/yugabyte \
--namespace nexustrade \
--set resource.master.requests.cpu=2 \
--set resource.master.requests.memory=4Gi \
--set resource.tserver.requests.cpu=4 \
--set resource.tserver.requests.memory=8Gi \
--set replicas.master=3 \
--set replicas.tserver=3 \
--set gflags.tserverysql_enable_auth=true \
--set gflags.masterysql_enable_auth=true
```

2. Initialize YugabyteDB Schema

```
bash
```

```
# Get YugabyteDB pod
YUGABYTE_POD=$(kubectl get pods -n nexustrade -l app=yb-tserver -o jsonpath='{.items[0].metadata.name}')

# Copy schema file
kubectl cp shared/prisma/schema.sql nexustrade/$YUGABYTE_POD:/tmp/

# Execute schema
kubectl exec -it -n nexustrade $YUGABYTE_POD --ysqlsh -h yb-tserver-0.yb-tservers.nexustrade -U yugabyte -d yugab
```

3. Deploy ScyllaDB Cluster

```
bash
```

```
# Add ScyllaDB operator
```

```
kubectl apply -f https://github.com/scylladb/scylla-operator/releases/download/v1.10.0/install.yaml
```

```
# Deploy ScyllaDB cluster
```

```
kubectl apply -f - <<EOF
```

```
apiVersion: scylla.scylladb.com/v1
```

```
kind: ScyllaCluster
```

```
metadata:
```

```
  name: nexustrade-scylla
```

```
  namespace: nexustrade
```

```
spec:
```

```
  version: 5.4.0
```

```
  datacenter:
```

```
    name: lagos
```

```
    racks:
```

```
      - name: rack1
```

```
        members: 3
```

```
        storage:
```

```
          capacity: 500Gi
```

```
          storageClassName: fast-ssd
```

```
        resources:
```

```
          requests:
```

```
            cpu: 4
```

```
            memory: 16Gi
```

```
          limits:
```

```
            cpu: 8
```

```
            memory: 32Gi
```

```
EOF
```

```
# Wait for cluster to be ready
```

```
kubectl wait --for=condition=Ready scyllacluster/nexustrade-scylla -n nexustrade --timeout=10m
```

4. Initialize ScyllaDB Schema

```
bash

# Get ScyllaDB pod
$SCYLLA_POD=$(kubectl get pods -n nexustrade -l app=scylla -o jsonpath='{.items[0].metadata.name}')

# Copy schema
kubectl cp shared/scylladb/schema.cql nexustrade/$SCYLLA_POD:/tmp/

# Execute schema
kubectl exec -it -n nexustrade $SCYLLA_POD -- cqlsh -f /tmp/schema.cql
```

5. Deploy DragonflyDB

```
bash

helm install dragonfly oci://ghcr.io/dragonflydb/dragonfly/helm/dragonfly \
--namespace nexustrade \
--set replicas=3 \
--set resources.requests.memory=4Gi \
--set resources.requests.cpu=2 \
--set persistence.enabled=true \
--set persistence.size=50Gi
```

6. Deploy MinIO

```
bash

helm install minio bitnami/minio \
--namespace nexustrade \
--set mode=distributed \
--set statefulset.replicaCount=4 \
--set persistence.size=500Gi \
--set auth.rootUser=nexustrade \
--set auth.rootPassword=$(openssl rand -base64 32)
```

7. Deploy Kafka

```
bash

helm install kafka bitnami/kafka \
--namespace nexustrade \
--set replicaCount=3 \
--set persistence.size=100Gi \
--set zookeeper.enabled=true \
--set zookeeper.replicaCount=3
```

Service Deployment

1. Create Secrets

```
bash

# Generate secrets
export JWT_SECRET=$(openssl rand -base64 64)
export JWT_REFRESH_SECRET=$(openssl rand -base64 64)
export DATABASE_PASSWORD=$(openssl rand -base64 32)

# Create Kubernetes secrets
kubectl create secret generic nexustrade-secrets \
--from-literal=database-url="postgresql://nexustrade:${DATABASE_PASSWORD}@yugabytedb:5433/nexustrade" \
--from-literal=jwt-secret="$JWT_SECRET" \
--from-literal=jwt-refresh-secret="$JWT_REFRESH_SECRET" \
--namespace nexustrade

# Store in Vault
vault kv put secret/nexustrade/production \
database-url="postgresql://nexustrade:${DATABASE_PASSWORD}@yugabytedb:5433/nexustrade" \
jwt-secret="$JWT_SECRET" \
jwt-refresh-secret="$JWT_REFRESH_SECRET"
```

2. Deploy Keycloak

```
bash
```

```
helm install keycloak bitnami/keycloak \
--namespace nexustrade \
--set auth.adminUser=admin \
--set auth.adminPassword=$(openssl rand -base64 32) \
--set production=true \
--set proxy=edge \
--set postgresql.enabled=true
```

3. Deploy HashiCorp Vault

```
bash
```

```
helm install vault hashicorp/vault \
--namespace nexustrade \
--set server.ha.enabled=true \
--set server.ha.replicas=3 \
--set ui.enabled=true

# Initialize Vault
kubectl exec -it vault-0 -n nexustrade -- vault operator init

# Unseal Vault (repeat for all pods)
kubectl exec -it vault-0 -n nexustrade -- vault operator unseal <unseal-key-1>
kubectl exec -it vault-0 -n nexustrade -- vault operator unseal <unseal-key-2>
kubectl exec -it vault-0 -n nexustrade -- vault operator unseal <unseal-key-3>
```

4. Deploy Real-Time Services

```
bash
```

```
# MongooseIM
kubectl apply -f infrastructure/kubernetes/mongooseim-deployment.yaml

# Janus Gateway
kubectl apply -f infrastructure/kubernetes/janus-deployment.yaml

# coturn
kubectl apply -f infrastructure/kubernetes/coturn-deployment.yaml

# Kamailio
kubectl apply -f infrastructure/kubernetes/kamailio-deployment.yaml
```

5. Deploy Microservices

```
bash

# API Gateway
kubectl apply -f infrastructure/kubernetes/api-gateway-deployment.yaml

# Trading Service
kubectl apply -f infrastructure/kubernetes/trading-deployment.yaml

# Bot Manager
kubectl apply -f infrastructure/kubernetes/bot-manager-deployment.yaml

# Bot Sidecar (Python)
kubectl apply -f infrastructure/kubernetes/bot-sidecar-deployment.yaml
```

6. Verify Deployments

```
bash

# Check all pods
kubectl get pods -n nexustrade

# Check services
kubectl get svc -n nexustrade

# Check ingress
kubectl get ingress -n nexustrade

# View logs
kubectl logs -f deployment/api-gateway -n nexustrade
```

Network Configuration

1. Configure Ingress

```

bash

kubectl apply -f - <<EOF
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: nexustrade-ingress
  namespace: nexustrade
  annotations:
    cert-manager.io/cluster-issuer: letsencrypt-prod
    nginx.ingress.kubernetes.io/ssl-redirect: "true"
    nginx.ingress.kubernetes.io/force-ssl-redirect: "true"
spec:
  ingressClassName: nginx
  tls:
    - hosts:
        - api.nexustrade.com
      secretName: nexustrade-tls
  rules:
    - host: api.nexustrade.com
      http:
        paths:
          - path: /
            pathType: Prefix
            backend:
              service:
                name: api-gateway
                port:
                  number: 80
EOF

```

2. Configure DNS

```

bash

# Get LoadBalancer IP
LB_IP=$(kubectl get svc ingress-nginx-controller -n ingress-nginx -o jsonpath='{.status.loadBalancer.ingress[0].ip}')

# Update DNS records
# - api.nexustrade.com → $LB_IP
# - ws.nexustrade.com → $LB_IP
# - turn.nexustrade.com → TURN server IP

```

3. Configure Geo-DNS

Use your DNS provider's Geo-DNS features to route traffic:

- Africa → Lagos/Johannesburg
 - Europe → London
 - Americas → Ashburn
 - Asia Pacific → Singapore
-

Security Hardening

1. Network Policies

```
bash  
# Apply network policies  
kubectl apply -f infrastructure/kubernetes/network-policies/
```

2. Pod Security Policies

```
bash

kubectl apply -f - <<EOF
apiVersion: policy/v1beta1
kind: PodSecurityPolicy
metadata:
  name: restricted
spec:
  privileged: false
  allowPrivilegeEscalation: false
  requiredDropCapabilities:
    - ALL
  volumes:
    - 'configMap'
    - 'emptyDir'
    - 'projected'
    - 'secret'
    - 'downwardAPI'
    - 'persistentVolumeClaim'
  runAsUser:
    rule: 'MustRunAsNonRoot'
  seLinux:
    rule: 'RunAsAny'
  fsGroup:
    rule: 'RunAsAny'
EOF
```

3. Enable mTLS

```
bash

# Install Istio service mesh (optional but recommended)
istioctl install --set profile=production

# Enable mTLS
kubectl apply -f - <<EOF
apiVersion: security.istio.io/v1beta1
kind: PeerAuthentication
metadata:
  name: default
  namespace: nexustrade
spec:
  mtls:
    mode: STRICT
EOF
```

Monitoring Setup

1. Deploy Prometheus

```
bash

helm install prometheus prometheus-community/kube-prometheus-stack \
--namespace monitoring \
--create-namespace \
--set prometheus.prometheusSpec.retention=30d \
--set prometheus.prometheusSpec.storageSpec.volumeClaimTemplate.spec.resources.requests.storage=100Gi
```

2. Deploy Grafana

```
bash

# Already included in kube-prometheus-stack
# Get Grafana password
kubectl get secret -n monitoring prometheus-grafana -o jsonpath="{.data.admin-password}" | base64 --decode
```

3. Deploy Loki

```
bash
```

```
helm install loki grafana/loki-stack \  
--namespace monitoring \  
--set loki.persistence.enabled=true \  
--set loki.persistence.size=100Gi
```

4. Deploy Tempo

```
bash
```

```
helm install tempo grafana/tempo \  
--namespace monitoring \  
--set tempo.persistence.enabled=true \  
--set tempo.persistence.size=50Gi
```

Post-Deployment Verification

1. Health Checks

```
bash
```

```
# API Gateway health  
curl https://api.nexustrade.com/health
```

```
# Database connectivity  
kubectl exec -it -n nexustrade deployment/api-gateway -- npm run db:check
```

```
# Redis connectivity  
kubectl exec -it -n nexustrade deployment/api-gateway -- npm run redis:check
```

2. Smoke Tests

```
bash
```

```
# Run smoke tests  
kubectl apply -f tests/smoke-tests.yaml
```

```
# View results  
kubectl logs -f job/smoke-tests -n nexustrade
```

3. Load Testing

```
bash

# Run load tests
k6 run tests/load-test.js
```

Rollback Procedures

Option 1: Helm Rollback

```
bash

# List releases
helm list -n nexustrade

# Rollback to previous version
helm rollback api-gateway -n nexustrade

# Rollback to specific revision
helm rollback api-gateway 3 -n nexustrade
```

Option 2: Kubernetes Rollback

```
bash

# View deployment history
kubectl rollout history deployment/api-gateway -n nexustrade

# Rollback to previous version
kubectl rollout undo deployment/api-gateway -n nexustrade

# Rollback to specific revision
kubectl rollout undo deployment/api-gateway --to-revision=2 -n nexustrade
```

Troubleshooting

Common Issues

1. Pods in CrashLoopBackOff

```
bash

# View logs
kubectl logs -f <pod-name> -n nexustrade

# Describe pod
kubectl describe pod <pod-name> -n nexustrade

# Check events
kubectl get events -n nexustrade --sort-by=.lastTimestamp'
```

2. Service Not Accessible

```
bash

# Check service
kubectl get svc <service-name> -n nexustrade

# Test from within cluster
kubectl run -it --rm debug --image=busybox --restart=Never -n nexustrade -- wget -O- http://<service-name>:port
```

3. Database Connection Issues

```
bash

# Test connection
kubectl exec -it -n nexustrade <pod-name> -- psql -h yugabytedb -U nexustrade -d nexustrade -c "SELECT 1"
```

4. High Memory Usage

```
bash

# Check resource usage
kubectl top pods -n nexustrade

# Increase resource limits
kubectl patch deployment api-gateway -n nexustrade -p '{"spec": {"template": {"spec": {"containers": [{"name": "api-gatewa
```

Support & Escalation

Support Channels

- **Critical (P0):** Page on-call engineer
- **High (P1):** Email: ops@nexustrade.com
- **Medium (P2):** Slack: #nexustrade-ops
- **Low (P3):** Ticket: support.nexustrade.com

Runbooks

- Database failover: [\(docs/runbooks/database-failover.md\)](#)
 - Service degradation: [\(docs/runbooks/service-degradation.md\)](#)
 - Security incident: [\(docs/runbooks/security-incident.md\)](#)
-

Deployment completed successfully! 

Monitor the system for 24 hours before declaring production-ready.